

Supervised Machine Learning: Regression

Peer To Peer Assignment

Main objective: Car Price Prediction

We are required to model the price of cars with the available independent variables. It will be used by the management to understand how exactly the prices vary with the independent variables. They can accordingly manipulate the design of the cars, the business strategy etc. to meet certain price levels.

Brief description of the data set and a summary of its attributes

I have data set consisting of 205 data points and 26 columns representing features.

There are some features that are not important for target variable , I will simply drop them and there are also some categorical feature, that I will convert to numerical form by Encoding method.

car_ID	ID of every car
CarName	Name of Car
fueltype	Type of Fuel
doornumber	Total number of door
carbody	Body of Car weather Sedan or Hatchback etc
enginelocation	Location of engine in car
wheelbase	Distance between rear and front wheel
carlength	Length of Car
carwidth	Width of Car
carheight	Height of Car
curbweight	Weight of Car without any passenger or item
enginetype	Type of engine
cylindernumber	Total cylinder in Car
enginesize	Size of Engine

boreratio Combustion Performance of Lean Burn Heavy-Duty Gaseous Engine

stroke A phase of the engine's cycle

horsepower Power of Car

peakrpm Revolution per minute

citympg City mileage per gallon

highwaympg Highway mileage per gallon

price Price of car

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df = pd.read_csv('E:\Machine Learning Course\Course2\week3\CarPrice.csv')
```

```
df.head()
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	...	enginesize	fuelsystem	boreratio	s
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	
4	5	2	audi 100ls	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	

5 rows × 26 columns

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column              Non-Null Count  Dtype
---  -
0   car_ID              205 non-null   int64
1   symboling           205 non-null   int64
2   CarName             205 non-null   object
3   fueltype            205 non-null   object
4   aspiration          205 non-null   object
5   doornumber          205 non-null   object
6   carbody             205 non-null   object
7   drivewheel         205 non-null   object
8   enginelocation      205 non-null   object
9   wheelbase          205 non-null   float64
10  carlength           205 non-null   float64
11  carwidth            205 non-null   float64
12  carheight           205 non-null   float64
13  curbweight          205 non-null   int64
14  enginetype          205 non-null   object
15  cylindernumber      205 non-null   object
16  enginesize          205 non-null   int64
17  fuelsystem          205 non-null   object
18  boreratio           205 non-null   float64
19  stroke              205 non-null   float64
20  compressionratio    205 non-null   float64
21  horsepower          205 non-null   int64
22  peakrpm             205 non-null   int64
```

```
df.isnull().sum()
```

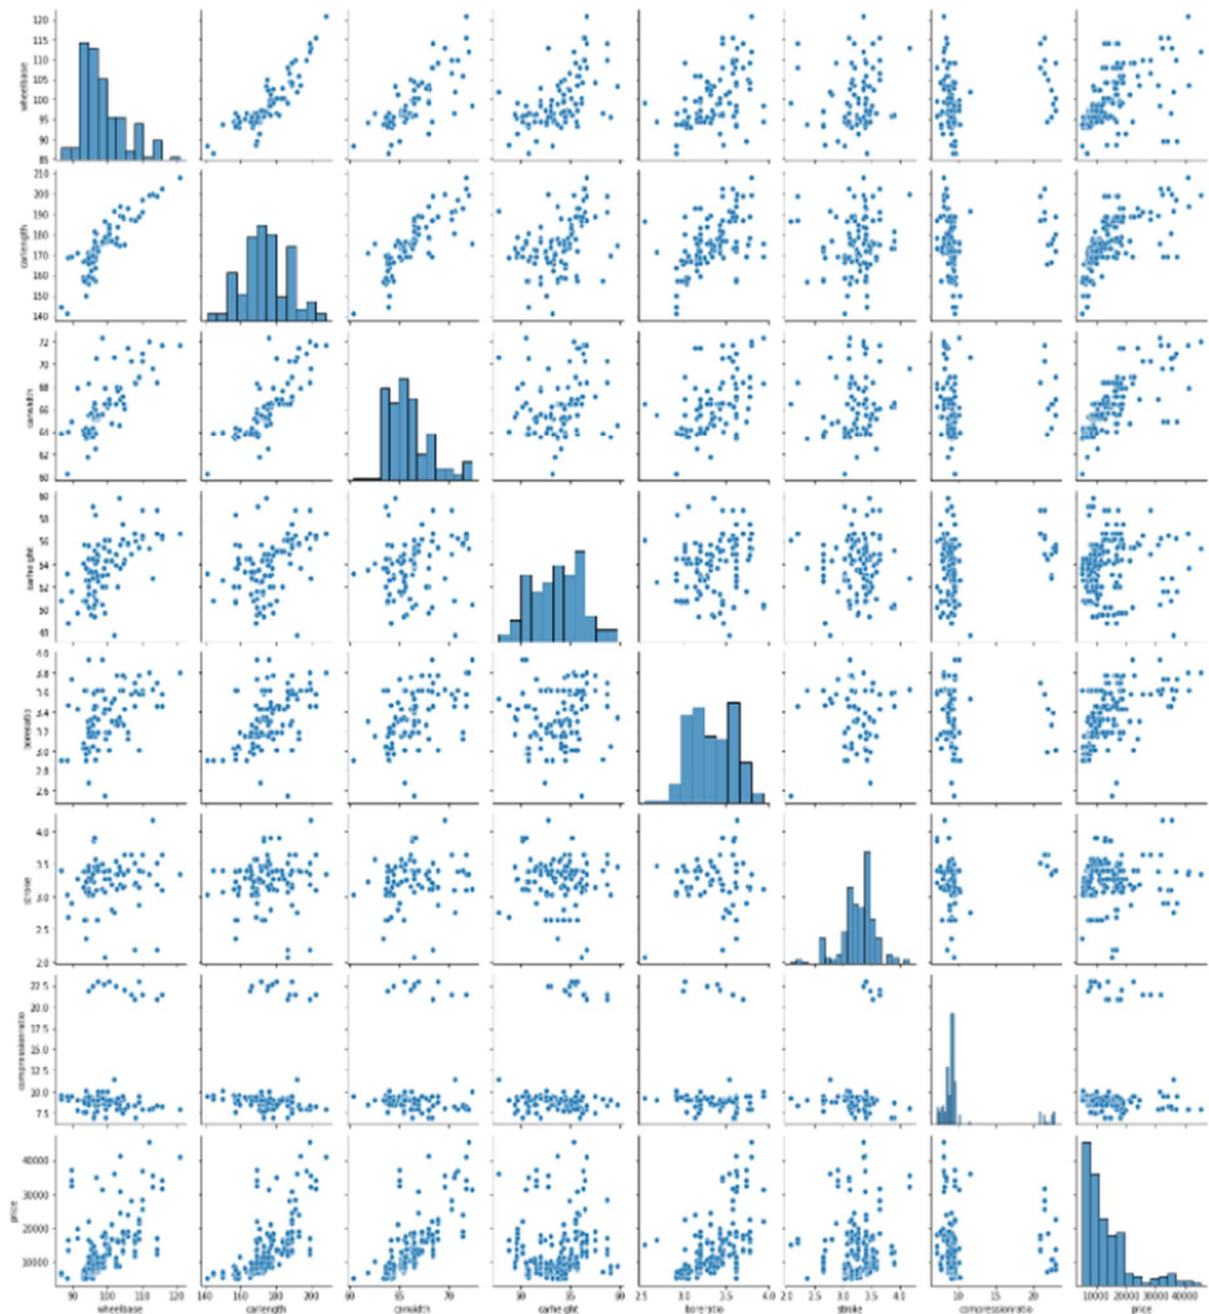
```
car_ID          0
symboling       0
CarName         0
fueltype        0
aspiration      0
doornumber      0
carbody         0
drivewheel      0
enginelocation  0
wheelbase       0
carlength       0
carwidth        0
carheight       0
curbweight      0
enginetype      0
cylindernumber  0
enginesize      0
fuelsystem      0
boreratio       0
stroke          0
compressionratio 0
horsepower      0
peakrpm         0
citympg         0
highwaympg      0
price           0
dtype: int64
```

```
df.drop(['fuelsystem', 'cylindernumber', 'enginetype', 'aspiration', 'symboling', 'car_ID'], inplace=True, axis=1)
```

```
small_df = df.loc[:, ['wheelbase', 'carlength', 'carwidth', 'carheight', 'bore:ratio', 'stroke', 'compressionratio', 'price']]
```

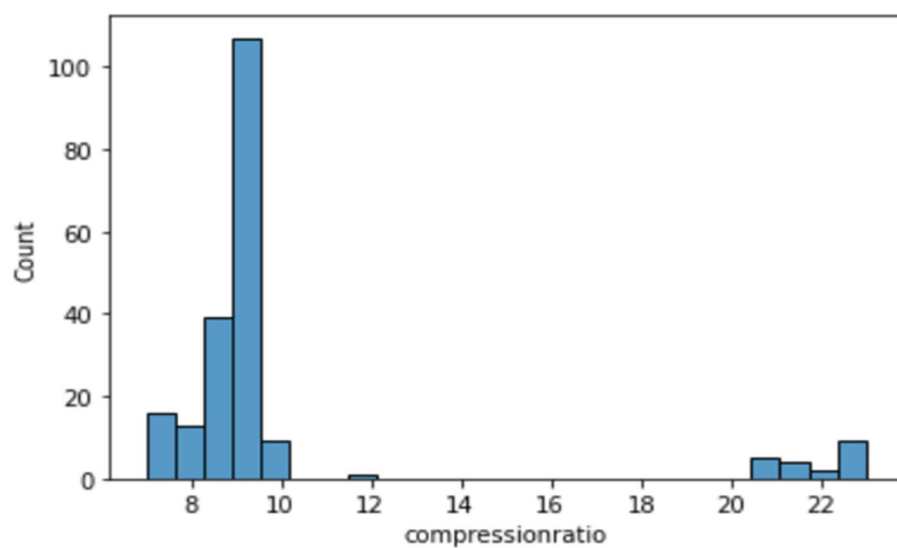
```
sns.pairplot(small_df)
```

<seaborn.axisgrid.PairGrid at 0x23446494b80>




```
sns.histplot(df['compressionratio'],bins=25)
```

```
<AxesSubplot:xlabel='compressionratio', ylabel='Count'>
```



```
sns.histplot(df['wheelbase'],bins=25)
```

```
<AxesSubplot:xlabel='wheelbase', ylabel='Count'>
```



```
mask = df.dtypes == np.float
float_cols = df.columns[mask]
```

```
skew_limit = 0.75 # define a limit above which we will log transform
skew_vals = df[float_cols].skew()
skew_vals
```

<ipython-input-141-30ffa34c1c>:1: DeprecationWarning: `np.float` is a deprecated alias for the builtin `float`. To silence this warning, use `float` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.float64` here.

Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
mask = df.dtypes == np.float
```


```
wheelbase      1.050214
carlength      0.155954
carwidth       0.904003
carheight      0.063123
boretostroke   0.020156
stroke        -0.689705
compressionratio 2.610862
price          1.777678
dtype: float64
```

```
skew_cols = (skew_vals.sort_values(ascending=False).to_frame().rename(columns={0: 'Skew'})).query('abs(Skew) > {}'.format(skew_limit))
skew_cols
```

Skew	
compressionratio	2.610862
price	1.777678
wheelbase	1.050214
carwidth	0.904003

```
for col in skew_cols.index.values:
    if col == "price":
        continue
    df[col] = df[col].apply(np.log1p)
```

```
plt.figure(figsize=(40,40))
sns.heatmap(df.corr(method='pearson'), cmap='PuBu', annot=True, linewidths=0.5)
plt.show()
```



```
feature = df.dtypes == np.object
```

```
<ipython-input-132-3fa00075be2d>:1: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is safe. Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
feature = df.dtypes == np.object
```

```
feature = df.dtypes[df.dtypes == np.object]
feature = feature.index.tolist()
```

```
<ipython-input-133-82993708c8bb>:1: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is safe. Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
feature = df.dtypes[df.dtypes == np.object]
```

```
df_dum = pd.get_dummies(df, columns=feature, drop_first=True)
```

```
df = df.drop(df[feature], axis=1)
```

```
pd.concat([df, df_dum], axis=1)
```

Summary of data exploration and Actions taken for Data Cleaning and Feature Engineering

1) FINDINGS

After reading data by pandas read_csv function, I applied

- isnull()** to see if there are any missing in the dataset. But I found there are no missing values in my dataset.
- Info()** function to see data type of different features of my dataset and I found there are many features that are important for my target variable and their data type is object.
- Skew() and hist()** function to check skewness of data and there were some columns that are right or left skewed. e.g. compressionratio and wheel etc.
- Pairplot()** to see correlation and also to find if there is any need to use **Polynomial feature** for higher degree. i.e. 2, 3, 4 etc.
- Heatmap()** to see correlation in more depth by printing their corresponding values of relation with each other.

2) ACTION:

- For column that are important for by target variable and their data type was Object, I applied **get_dummies()** function of pandas to convert them into numeric type.
- For removing skewness of different columns, I applied **log1p** transformation function.
- Although I found there are some columns are correlated with each other but when I applied **Variation Inflation Factor** technique to remove correlation among them, I found that it have negative impact on my **r2_score** because in this technique some columns are drop for eliminating correlation. So then I trained my model without removing correlation.
- Also their was some columns that are not important for our target, So I simply drop them.
- Similarly, I used **Standard Scaler** and **Polynomial feature** technique on dataset before giving it to model for training.

Summary of Three Different Linear Regression Model

1) Simple Linear Regression Model

```
X = df.drop('price',axis=1)
y = df['price']
```

```
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.model_selection import GridSearchCV, train_test_split, KFold
from sklearn.linear_model import LinearRegression,Lasso,Ridge
from sklearn.pipeline import Pipeline
from sklearn.metrics import r2_score
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.25,random_state=72018)
```

```
s = StandardScaler()
X_train_s = s.fit_transform(X_train)
X_test_s = s.transform(X_test)
```

```
lr = LinearRegression()
lr = lr.fit(X_train_s,y_train)
pred = lr.predict(X_test_s)
r2_score(pred,y_test)
```

```
0.6009998432027648
```

2) Ridge Regression Model

```
kf = KFold(shuffle=True,random_state=72018,n_splits=3)
```

```
estimator = Pipeline([
    ("Polynomail_feature",PolynomialFeatures()),
    ("scaler",StandardScaler()),
    ("ridge_regression",Ridge())
])

param = {
    'Polynomail_feature__degree':[1,2,3],
    'ridge_regression__alpha':np.geomspace(4,20,30)
}

grid = GridSearchCV(estimator,param,cv=kf)
```

```
grid.fit(X,y)
```

```
GridSearchCV(cv=KFold(n_splits=3, random_state=72018, shuffle=True),
             estimator=Pipeline(steps=[('Polynomail_feature',
                                         PolynomialFeatures()),
                                         ('scaler', StandardScaler()),
                                         ('ridge_regression', Ridge())]),
             param_grid={'Polynomail_feature__degree': [1, 2, 3],
                         'ridge_regression__alpha': array([ 4.
5.27924796,  5.58051751,  5.89897953,  6.23561514,  6.59146146,
6.96761476,  7.36523392,  7.78554391,  8.22983963,  8.69948987,
9.19594151,  9.72072404, 10.27545421, 10.86184103, 11.48169104,
12.13691388, 12.82952815, 13.56166768, 14.33558803, 15.15367351,
16.01844446, 16.93256509, 17.89885162, 18.92028098, 20.
])})
```

```
pred_r = grid.predict(X)
```

```
r2_score(pred_r,y)
```

```
0.8940782687472872
```

```
grid.best_score_, grid.best_params_
```

```
(0.8229369717953522,
 {'Polynomail_feature__degree': 3,
  'ridge_regression__alpha': 12.82952815374728})
```

3) Lasso Regression Model

```

estimator = Pipeline([
    ("Polynomail_feature",PolynomialFeatures()),
    ("scaler",StandardScaler()),
    ("lasso_regression",Lasso(max_iter=100000))
])

param = {
    'Polynomail_feature__degree':[1,2,3],
    'lasso_regression__alpha':np.geomspace(0.001,10,5)
}

grid_1 = GridSearchCV(estimator,param,cv=kf)

```

```
grid_1.fit(X,y)
```

```

GridSearchCV(cv=KFold(n_splits=3, random_state=72018, shuffle=True),
             estimator=Pipeline(steps=[('Polynomail_feature',
                                         PolynomialFeatures()),
                                         ('scaler', StandardScaler()),
                                         ('lasso_regression',
                                          Lasso(max_iter=100000))])),
             param_grid={'Polynomail_feature__degree': [1, 2, 3],
                         'lasso_regression__alpha': array([1.e-03, 1.e-02, 1.e-

```

```
pred_1 = grid_1.predict(X)
```

```
r2_score(pred_1,y)
```

```
0.895975904352647
```

```
grid_1.best_score_, grid_1.best_params_
```

```

(0.8177757224961648,
 {'Polynomail_feature__degree': 2, 'lasso_regression__alpha': 10.0})

```

1) As in first case I applied simple linear regression on model and it's r2_score was 0.60 which is not considered as good result.

2) In second case I applied Polynomial feature technique on dataset before giving it to Ridge regression. I also used GridSearchCV method and pipeline technique to make process fast and easy to find out best best parameters for our model to have good prediction.

So by doing this, I found that parameters i.e polynomial feature and alpha values 3 and 14 are best for our model to have r^2_score of 0.89 .

3) Similarly for case three I used Lasso Regression, GridSearchCV and pipeline technique and founded that value of 2 for degree and 10 for alpha is best for our model to have good r^2_score of 0.89 .

A paragraph explaining which of your regressions you recommend as a final model that best fits your needs in terms of accuracy and explainability.

For choosing model that best fir our data, Ridge and Lasso both are better than simple linear regression and both have same r^2_score , So we can choose any one from both of them but if we want interpretability along with our main goal of prediction then Lasso will be choice.

Summary Key Findings and Insights, which walks your reader through the main drivers of your model and insights from your data derived from your linear regression model

Some of key point for this data and regression model is that as we know that correlation in data is not good for our model and it have impact on model accuracy but in our case when we used technique of VIF to eliminate correlation from data, I found that it does not have good impact on accuracy of model because in this technique we drop one of column that are correlated with each other, So by doing that we will end up with having less number of feature for our model to train on which is again problem. i.e Problem of underfitting.

Suggestions for next steps in analyzing this data, which may include suggesting revisiting this model adding specific data features to achieve a better explanation or a better prediction

So my next suggestion for analyzing this data will be to have more data because as we can see our data consist of only 205 row or observation that is not good enough to train model.
