

# Bank Customer Churn

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import os
```

## Objective:

*Churn prevention allows companies to develop loyalty programs and retention campaigns to keep as many customers as possible.*

## Dataset Description:

RowNumber—corresponds to the record (row) number and has no effect on the output.

CustomerId—contains random values and has no effect on customer leaving the bank.

Surname—the surname of a customer has no impact on their decision to leave the bank.

CreditScore—can have an effect on customer churn, since a customer with a higher credit score is less likely to leave the bank.

Geography—a customer's location can affect their decision to leave the bank.

Gender—it's interesting to explore whether gender plays a role in a customer leaving the bank.

Age—this is certainly relevant, since older customers are less likely to leave their bank than younger ones.

Tenure—refers to the number of years that the customer has been a client of the bank. Normally, older clients are more loyal and less likely to leave a bank.

Balance—also a very good indicator of customer churn, as people with a higher balance in their accounts are less likely to leave the bank compared to those with lower balances.

NumOfProducts—refers to the number of products that a customer has purchased through the bank.

HasCrCard—denotes whether or not a customer has a credit card. This column is also relevant, since people with a credit card are less likely to leave the bank.

IsActiveMember—active customers are less likely to leave the bank.

EstimatedSalary—as with balance, people with lower salaries are more likely to leave the bank compared to those with higher salaries.

Exited—whether or not the customer left the bank.

```
]:
```

```
]:
```

```
df = pd.read_csv('churn.csv')
df.head()
```

```
df = pd.read_csv('churn.csv')
df.head()
```

```
]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	1013
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	1125
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	1139
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	938
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	790

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   RowNumber       10000 non-null  int64
1   CustomerId      10000 non-null  int64
2   Surname         10000 non-null  object
3   CreditScore     10000 non-null  int64
4   Geography       10000 non-null  object
5   Gender         10000 non-null  object
6   Age            10000 non-null  int64
7   Tenure         10000 non-null  int64
```

```
7   Tenure         10000 non-null  int64
8   Balance        10000 non-null  float64
9   NumOfProducts  10000 non-null  int64
10  HasCrCard      10000 non-null  int64
11  IsActiveMember 10000 non-null  int64
12  EstimatedSalary 10000 non-null  float64
13  Exited         10000 non-null  int64
```

```
dtypes: float64(2), int64(9), object(3)
```

```
memory usage: 1.1+ MB
```

```
df.shape
```

```
:(10000, 14)
```

```
df.dtypes
```

```
: RowNumber      int64
  CustomerId     int64
  Surname        object
  CreditScore    int64
  Geography      object
  Gender         object
  Age           int64
  Tenure        int64
  Balance       float64
  NumOfProducts int64
  HasCrCard     int64
  IsActiveMember int64
```

```
IsActiveMember      int64
EstimatedSalary     float64
Exited              int64
dtype: object
```

```
df.isnull().sum()
```

```
] RowNumber      0
   CustomerId    0
   Surname       0
   CreditScore    0
   Geography     0
   Gender        0
   Age          0
   Tenure       0
   Balance      0
   NumOfProducts 0
   HasCrCard    0
   IsActiveMember 0
   EstimatedSalary 0
   Exited       0
dtype: int64
```

```
df.drop(['RowNumber', 'CustomerId', 'Surname'], axis=1, inplace=True)
```

```
df.shape
```

```
] (10000, 11)
```

```
df.Geography.unique()
```

```
] array(['France', 'Spain', 'Germany'], dtype=object)
```

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
object_col = df.dtypes[df.dtypes==np.object]
object_col = object_col.index.to_list()
```

```
object_col
```

```
] ['Geography', 'Gender']
```

```
for col in object_col:
    df[col] = le.fit_transform(df[col])
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   CreditScore  10000 non-null  int64
1   Geography   10000 non-null  int32
2   Gender      10000 non-null  int32
3   Age         10000 non-null  int64
4   Tenure      10000 non-null  int64
```

```
4  Tenure          10000 non-null int64
5  Balance         10000 non-null float64
6  NumOfProducts  10000 non-null int64
7  HasCrCard       10000 non-null int64
8  IsActiveMember  10000 non-null int64
9  EstimatedSalary 10000 non-null float64
10 Exited          10000 non-null int64
dtypes: float64(2), int32(2), int64(7)
memory usage: 781.4 KB
```

```
df = df.astype('float')
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   CreditScore      10000 non-null  float64
1   Geography        10000 non-null  float64
2   Gender           10000 non-null  float64
3   Age              10000 non-null  float64
4   Tenure           10000 non-null  float64
5   Balance          10000 non-null  float64
6   NumOfProducts    10000 non-null  float64
7   HasCrCard        10000 non-null  float64
8   IsActiveMember   10000 non-null  float64
9   EstimatedSalary  10000 non-null  float64
10  Exited           10000 non-null  float64
```

```
9   EstimatedSalary 10000 non-null float64
10  Exited          10000 non-null float64
dtypes: float64(11)
memory usage: 859.5 KB
```

```
feature_col = [x for x in df.columns if x!='Exited']
```

```
df.describe()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	650.528800	0.746300	0.545700	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.2398
std	96.653299	0.827529	0.497932	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.4928
min	350.000000	0.000000	0.000000	18.000000	0.000000	0.000000	1.000000	0.000000	0.000000	11.5800
25%	584.000000	0.000000	0.000000	32.000000	3.000000	0.000000	1.000000	0.000000	0.000000	51002.1100
50%	652.000000	0.000000	1.000000	37.000000	5.000000	97198.540000	1.000000	1.000000	1.000000	100193.9150
75%	718.000000	1.000000	1.000000	44.000000	7.000000	127644.240000	2.000000	1.000000	1.000000	149388.2475
max	850.000000	2.000000	1.000000	92.000000	10.000000	250898.090000	4.000000	1.000000	1.000000	199992.4800

```
df_org = df.copy()
corr = df.corr()
corr
```

```
df_org = df.copy()
corr = df.corr()
corr
```

```
]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
CreditScore	1.000000	0.007888	-0.002857	-0.003965	0.000842	0.006268	0.012238	-0.005458	0.025651	-0.001384	-0.027094
Geography	0.007888	1.000000	0.004719	0.022812	0.003739	0.069408	0.003972	-0.008523	0.006724	-0.001369	0.035943
Gender	-0.002857	0.004719	1.000000	-0.027544	0.014733	0.012087	-0.021859	0.005766	0.022544	-0.008112	-0.106512
Age	-0.003965	0.022812	-0.027544	1.000000	-0.009997	0.028308	-0.030680	-0.011721	0.085472	-0.007201	0.285323
Tenure	0.000842	0.003739	0.014733	-0.009997	1.000000	-0.012254	0.013444	0.022583	-0.028362	0.007784	-0.014001
Balance	0.006268	0.069408	0.012087	0.028308	-0.012254	1.000000	-0.304180	-0.014858	-0.010084	0.012797	0.118533
NumOfProducts	0.012238	0.003972	-0.021859	-0.030680	0.013444	-0.304180	1.000000	0.003183	0.009612	0.014204	-0.047820
HasCrCard	-0.005458	-0.008523	0.005766	-0.011721	0.022583	-0.014858	0.003183	1.000000	-0.011866	-0.009933	-0.007138
IsActiveMember	0.025651	0.006724	0.022544	0.085472	-0.028362	-0.010084	0.009612	-0.011866	1.000000	-0.011421	-0.156128
EstimatedSalary	-0.001384	-0.001369	-0.008112	-0.007201	0.007784	0.012797	0.014204	-0.009933	-0.011421	1.000000	0.012097
Exited	-0.027094	0.035943	-0.106512	0.285323	-0.014001	0.118533	-0.047820	-0.007138	-0.156128	0.012097	1.000000

```
for x in range(corr.shape[0]):
    corr.iloc[x,x] = 0.0
corr
```

```
for x in range(corr.shape[0]):
    corr.iloc[x,x] = 0.0
corr
```

```
]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
CreditScore	0.000000	0.007888	-0.002857	-0.003965	0.000842	0.006268	0.012238	-0.005458	0.025651	-0.001384	-0.027094
Geography	0.007888	0.000000	0.004719	0.022812	0.003739	0.069408	0.003972	-0.008523	0.006724	-0.001369	0.035943
Gender	-0.002857	0.004719	0.000000	-0.027544	0.014733	0.012087	-0.021859	0.005766	0.022544	-0.008112	-0.106512
Age	-0.003965	0.022812	-0.027544	0.000000	-0.009997	0.028308	-0.030680	-0.011721	0.085472	-0.007201	0.285323
Tenure	0.000842	0.003739	0.014733	-0.009997	0.000000	-0.012254	0.013444	0.022583	-0.028362	0.007784	-0.014001
Balance	0.006268	0.069408	0.012087	0.028308	-0.012254	0.000000	-0.304180	-0.014858	-0.010084	0.012797	0.118533
NumOfProducts	0.012238	0.003972	-0.021859	-0.030680	0.013444	-0.304180	0.000000	0.003183	0.009612	0.014204	-0.047820
HasCrCard	-0.005458	-0.008523	0.005766	-0.011721	0.022583	-0.014858	0.003183	0.000000	-0.011866	-0.009933	-0.007138
IsActiveMember	0.025651	0.006724	0.022544	0.085472	-0.028362	-0.010084	0.009612	-0.011866	0.000000	-0.011421	-0.156128
EstimatedSalary	-0.001384	-0.001369	-0.008112	-0.007201	0.007784	0.012797	0.014204	-0.009933	-0.011421	0.000000	0.012097
Exited	-0.027094	0.035943	-0.106512	0.285323	-0.014001	0.118533	-0.047820	-0.007138	-0.156128	0.012097	0.000000

```
corr.abs().idxmax()
```

```
]: CreditScore      Exited
   Geography      Balance
```

```

Gender          Exited
Age             Exited
Tenure          IsActiveMember
Balance         NumOfProducts
NumOfProducts  Balance
HasCrCard       Tenure
IsActiveMember  Exited
EstimatedSalary NumOfProducts
Exited         Age
dtype: object

```

```

skew_col = df.skew().sort_values(ascending=False)
if 'Exited' in skew_col:
    skew_col.drop('Exited',inplace=True)

```

```

skew_col = skew_col.loc[skew_col>0.75]
skew_col

```

```

Age    1.01132
dtype: float64

```

```

for col in skew_col.index.to_list():
    df[col] = np.log1p(df[col])

```

```

skew_col = df.skew().sort_values(ascending=False)
skew_col

```

```

Exited          1.471611
NumOfProducts   0.745568
Geography       0.500916
Age             0.203360
Tenure          0.010991
EstimatedSalary 0.002085
IsActiveMember  -0.060437


```

```

ax = plt.figure(figsize=(12,8))
sns.heatmap(corr,annot=True)

```

<AxesSubplot:>



```
from sklearn.preprocessing import StandardScaler
ss = StandardScaler()
df[feature_col] = ss.fit_transform(df[feature_col])
```

## Summary of EDA:

After loading dataset i checked for null values but their was no missing in any column.

Then i checked skewness and applied log1p transform to make it like normal disturbution.

Checked Correlation but didn't tooked any action beacuse dataset donot have alot of features to drop them for elimation of correlation.

Dropped Unnecessary columns from dataset.

As we know its not good idea to cluster dataset without scaling data, So before applying clustering technique i scaleed data.

## Different Models Trained Using Different Clustering Algorithms

```
from sklearn.cluster import KMeans

km_list = list()

for i in range(1,15):
    km = KMeans(n_clusters=i,random_state=42)
    km = km.fit(df[feature_col])


    km_list.append(pd.Series({'clusters':i,'inertia':km.inertia_, 'model':km}))
```

```
plot_df = pd.concat(km_list,axis=1).T[['clusters','inertia']].set_index('clusters')

ax = plot_df.plot(marker='o',ls='--')
ax.set_xticks(range(0,15,2))
```

```
ax.set_xlim(0,15)
ax.set(xlabel='clusters',ylabel='inertia')
```

```
[Text(0.5, 0, 'clusters'), Text(0, 0.5, 'inertia')]
```



```
from sklearn.cluster import KMeans
km = KMeans(n_clusters=2,random_state=42)
km = km.fit(df[feature_col])

df['kmeans'] = km.predict(df[feature_col])
```

```
df.head(4)
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited	kmeans
0	-0.326221	-0.901886	-1.095988	0.423222	-1.041760	-1.225848	-0.911583	0.646092	0.970243	0.021886	1.0	0
1	-0.440036	1.515067	-1.095988	0.329713	-1.387538	0.117350	-0.911583	-1.547768	0.970243	0.216534	0.0	1



1	-0.440036	1.515067	-1.095988	0.329713	-1.387538	0.117350	-0.911583	-1.547768	0.970243	0.216534	0.0	1
2	-1.536794	-0.901886	-1.095988	0.423222	1.032908	1.333053	2.527057	0.646092	-1.030670	0.240687	1.0	0
3	0.501521	-0.901886	-1.095988	0.135823	-1.387538	-1.225848	0.807737	-1.547768	-1.030670	-0.108918	0.0	0

```
X = df.drop('Exited',axis=1)
y = df['Exited']
```

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score,classification_report
from sklearn.model_selection import StratifiedShuffleSplit
```

```
sss = StratifiedShuffleSplit(n_splits=5,random_state=42)

def get_avg_roc(estimator,X,y):
    roc_auc_list = []
    for train_idx,test_idx in sss.split(X,y):
        X_train,X_test = X.iloc[train_idx],X.iloc[test_idx]
        y_train,y_test = y.iloc[train_idx],y.iloc[test_idx]
        estimator.fit(X_train,y_train)
        y_pred = estimator.predict(X_test)
        y_scored = estimator.predict_proba(X_test)[:,:1]
        roc_auc_list.append(roc_auc_score(y_test,y_scored))
    return np.mean(roc_auc_list)

estimator = RandomForestClassifier()
roc_auc = get_avg_roc(estimator,X,y)
print("Using kmeans cluster as input to random forest, roc-auc is \"{}\".format(roc_auc))
```

Using kmeans cluster as input to random forest, roc-auc is "0.8476814218149572"

```
from sklearn.cluster import AgglomerativeClustering
```

```
from sklearn.cluster import AgglomerativeClustering

X2 = df[feature_col]
y = df['Exited']
```

```
ag = AgglomerativeClustering(n_clusters=2,linkage='ward',compute_full_tree=True)
ag = ag.fit(X2)
X2['agglo'] = ag.fit_predict(X2)
```

<ipython-input-243-4329a8be967b>:3: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead  
  
See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
X2['agglo'] = ag.fit\_predict(X2)

```
sss = StratifiedShuffleSplit(n_splits=5,random_state=42)

def get_avg_roc(estimator,X,y):
    roc_auc_list = []
    for train_idx,test_idx in sss.split(X,y):
        X_train,X_test = X.iloc[train_idx],X.iloc[test_idx]
        y_train,y_test = y.iloc[train_idx],y.iloc[test_idx]
        estimator.fit(X_train,y_train)
        y_pred = estimator.predict(X_test)
        y_scored = estimator.predict_proba(X_test)[:,:1]
        roc_auc_list.append(roc_auc_score(y_test,y_scored))
    return np.mean(roc_auc_list)

estimator = RandomForestClassifier()
roc_auc = get_avg_roc(estimator,X2,y)
print("Using Agglomerative cluster as input to random forest, roc-auc is \"{}\".format(roc_auc))
```

Using Agglomerative cluster as input to random forest, roc-auc is "0.8460685781850428"



```
from sklearn.cluster import DBSCAN
```

```
X3 = df[feature_col]  
y = df['Exited']
```

```
db = DBSCAN(eps=0.5,min_samples=3)  
db = db.fit(X3)  
X3['dbscan'] = db.fit_predict(X3)
```

<ipython-input-263-caba3fd15feb>:3: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
X3['dbscan'] = db.fit\_predict(X3)

```
sss = StratifiedShuffleSplit(n_splits=5,random_state=42)
```

```
def get_avg_roc(estimator,X,y):  
    roc_auc_list = []  
    for train_idx,test_idx in sss.split(X,y):  
        X_train,X_test = X.iloc[train_idx],X.iloc[test_idx]  
        y_train,y_test = y.iloc[train_idx],y.iloc[test_idx]  
        estimator.fit(X_train,y_train)  
        y_pred = estimator.predict(X_test)  
        y_scored = estimator.predict_proba(X_test)[:,:1]  
        roc_auc_list.append(roc_auc_score(y_test,y_scored))  
    return np.mean(roc_auc_list)
```

```
estimator = RandomForestClassifier()  
roc_auc = get_avg_roc(estimator,X3,y)  
print("Using DBSCAN cluster as input to random forest, roc-auc is {}".format(roc_auc))
```

Using DBSCAN cluster as input to random forest, roc-auc is "0.8497339639373337"

Using DBSCAN cluster as input to random forest, roc-auc is "0.8497339639373337"

## Summary and Explanation of Unsupervised Learning Models

I have applied different clustering techniques on this dataset and found that DBSCAN performs little bit better than other clustering technique. And during hyperparameter tuning I found that DBSCAN have high accuracy for value of number of sample equal to 3 instead of default value of 5.

## Summary

Actually in this dataset we have feature of exited or not but that was only for our testing purpose, we want to see how much our clustering technique combine with supervised learning is showing us accurate result. So as key note we can cluster our clients into churn and will not-churn, So that you make policies accordingly.

## Suggestion

Although we don't have complex data but still you can apply PCA technique to come up with new feature and see if there is any change in accuracy of model etc.